

# task\_1\_random\_forest

February 22, 2026

```
[1]: # # Mount Google Drive
# from google.colab import drive
# drive.mount('/content/drive')

# Imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# Scikit-Learn Models & Tools
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import classification_report, confusion_matrix,
    accuracy_score, f1_score, log_loss
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.base import BaseEstimator, TransformerMixin

# Ignore warnings for cleaner output
import warnings
warnings.filterwarnings('ignore')

print("Libraries imported successfully!")
```

Libraries imported successfully!

```
[2]: def load_and_prep_data():
    train_path = '/content/drive/MyDrive/University/CSCI316 - Big Data Mining
    ↵Techniques/Group Assignment/UNSW_NB15_training-set.csv'
    test_path = '/content/drive/MyDrive/University/CSCI316 - Big Data Mining
    ↵Techniques/Group Assignment/UNSW_NB15_testing-set.csv'

    # Local ENV file path
    test_path = '/Users/jju/Documents/SIM/Semester 1, 2026/CSCI316 Big Data
    ↵Mining/Assignments/Group_Assignment_Database/UNSW_NB15_testing-set.csv'
```

```

train_path = '/Users/jju/Documents/SIM/Semester 1, 2026/CSCI316 Big DataMining/Assignments/Group_Assignment_Database/UNSW_NB15_training-set.csv'

# 1. Load both datasets
print("Loading data...")
df_train_orig = pd.read_csv(train_path)
df_test_orig = pd.read_csv(test_path)

# 2. Combine them
df_full = pd.concat([df_train_orig, df_test_orig], axis=0).
    ↪reset_index(drop=True)

missing_data = df_full.isnull().sum()
if missing_data.sum() > 0:
    print("\n[WARNING] Missing values detected:")
    print(missing_data[missing_data > 0])
else:
    print("\n[CHECK] No true missing values (NaN) found.")

# 3. Define Features and Target
# We use 'label' for stratification to maintain class balance
drop_cols = ['label', 'id', 'attack_cat']
X = df_full.drop(drop_cols, axis=1)
y = df_full['label']

# 4. Stratified Split: Train (70%) and Temp (30%)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# 5. Split Temp into Validation (15%) and Test (15%)
# 0.5 * 30% = 15%
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)

# Now X_train, X_val, X_test/ y_train, y_val, y_test available

# 6. Preprocessing (Label Encoding)
le = LabelEncoder()
categorical_cols = ['proto', 'service', 'state'] # attack_cat removed since
    ↪already dropped
for col in categorical_cols:
    # labelEncode each train/test/validation data
    X_train[col] = le.fit_transform(X_train[col].astype(str))
    X_val[col] = le.fit_transform(X_val[col].astype(str))
    X_test[col] = le.fit_transform(X_test[col].astype(str))

```

```

print(f"Features used for training ({len(X.columns)} total):")
print(list(X.columns))
print("-" * 50)

print(f"Data Loaded and Split:")
print(f"Train: {X_train.shape}, Val: {X_val.shape}, Test: {X_test.shape}")

return X_train, y_train, X_val, y_val, X_test, y_test

# Unpack the validation set
X_train, y_train, X_val, y_val, X_test, y_test = load_and_prep_data()

def check_distributions(y_train, y_val, y_test):
    sets = {'Train': y_train, 'Validation': y_val, 'Test': y_test}
    for name, set_data in sets.items():
        counts = set_data.value_counts(normalize=True) * 100
        print(f"{name} Distribution: Normal: {counts[0]:.2f}%, Attack:{counts[1]:.2f}%")

check_distributions(y_train, y_val, y_test)

```

Loading data...

[CHECK] No true missing values (NaN) found.

Features used for training (42 total):

- ['dur', 'proto', 'service', 'state', 'spkts', 'dpkts', 'sbytes', 'dbytes', 'rate', 'sttl', 'dttl', 'sload', 'dload', 'sloss', 'dloss', 'sinpkt', 'dinpkt', 'sjit', 'djit', 'swin', 'stcpb', 'dtcpb', 'dwin', 'tcprrtt', 'synack', 'ackdat', 'smean', 'dmean', 'trans\_depth', 'response\_body\_len', 'ct\_srv\_src', 'ct\_state\_ttl', 'ct\_dst\_ltm', 'ct\_src\_dport\_ltm', 'ct\_dst\_sport\_ltm', 'ct\_dst\_src\_ltm', 'is\_ftp\_login', 'ct\_ftp\_cmd', 'ct\_flw\_http\_mthd', 'ct\_src\_ltm', 'ct\_srv\_dst', 'is\_sm\_ips\_ports']

-----

Data Loaded and Split:

Train: (180371, 42), Val: (38651, 42), Test: (38651, 42)

Train Distribution: Normal: 36.09%, Attack: 63.91%

Validation Distribution: Normal: 36.09%, Attack: 63.91%

Test Distribution: Normal: 36.09%, Attack: 63.91%

[3]:

```

class CustomFeatureTransformer(BaseEstimator, TransformerMixin):
    def __init__(self, use_extra_features=True):
        self.use_extra_features = use_extra_features

    def fit(self, X, y=None):

```

```

    return self

def transform(self, X):
    X_transformed = X.copy()
    if self.use_extra_features:
        if 'spkts' in X.columns and 'dpkts' in X.columns:
            X_transformed['pkt_ratio'] = (X_transformed['spkts'] + 1)/
            ↵(X_transformed['dpkts'] + 1)
        if 'sttl' in X.columns and 'dttl' in X.columns:
            X_transformed['ttl_gap'] = abs(X_transformed['sttl'] -
            ↵X_transformed['dttl'])
    return X_transformed

```

```

[4]: def train_and_tune_rf(X_train, y_train, X_val, y_val):
    print("Starting Grid Search for Random Forest...")
    start_time = time.time()

    pipeline = Pipeline([
        ('custom_transformer', CustomFeatureTransformer()),
        ('scaler', StandardScaler()),
        ('rf', RandomForestClassifier(random_state=42, n_jobs=-1))
    ])

    param_grid = {
        'custom_transformer__use_extra_features': [True, False],
        'rf__n_estimators': [50, 100],
        'rf__max_depth': [None, 10, 20],
        'rf__criterion': ['gini', 'entropy']
    }

    grid_search = GridSearchCV(
        pipeline,
        param_grid,
        cv=3,
        scoring='f1',
        verbose=1
    )

    grid_search.fit(X_train, y_train)
    results_df = pd.DataFrame(grid_search.cv_results_)

    # Extract and display all combination results
    results_df = pd.DataFrame(grid_search.cv_results_)
    relevant_results = results_df[['params', 'mean_test_score', ↵
    ↵'std_test_score', 'rank_test_score']]
    relevant_results = relevant_results.sort_values(by='rank_test_score')

```

```

print("\n--- All Hyperparameter Combinations (Sorted by F1-Score) ---")
print(relevant_results.to_string(index=False))

# Validate the best model on the Validation set
val_predictions = grid_search.best_estimator_.predict(X_val)
val_f1 = f1_score(y_val, val_predictions)

print(f"\nGrid Search Complete in {time.time() - start_time:.2f} seconds.")
print(f"Best Parameters: {grid_search.best_params_}")
print(f"Validation F1-Score: {val_f1:.4f}")

return grid_search.best_estimator_

# Run the training
best_rf_model = train_and_tune_rf(X_train, y_train, X_val, y_val)

```

Starting Grid Search for Random Forest...

Fitting 3 folds for each of 24 candidates, totalling 72 fits

	params	mean_test_score	std_test_score
rank_test_score			
1	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy', 'rf__max_depth': None, 'rf__n_estimators': 100}	0.958350	0.000867
2	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini', 'rf__max_depth': None, 'rf__n_estimators': 100}	0.958277	0.000280
3	{'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy', 'rf__max_depth': None, 'rf__n_estimators': 100}	0.958161	0.000639
4	{'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini', 'rf__max_depth': None, 'rf__n_estimators': 100}	0.958156	0.000565
5	{'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini', 'rf__max_depth': None, 'rf__n_estimators': 50}	0.958063	0.000621
6	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy', 'rf__max_depth': None, 'rf__n_estimators': 50}	0.958032	0.000764
7	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini', 'rf__max_depth': None, 'rf__n_estimators': 50}	0.957966	0.000378
8	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy', 'rf__max_depth': 20, 'rf__n_estimators': 100}	0.957827	0.000626
	{'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini', 'rf__max_depth': 20, 'rf__n_estimators': 100}	0.957762	0.000339

```

9      {'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy',
'rf__max_depth': 20, 'rf__n_estimators': 50}           0.957724      0.000719
10     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy',
'rf__max_depth': None, 'rf__n_estimators': 50}          0.957646      0.000512
11     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini',
'rf__max_depth': 20, 'rf__n_estimators': 100}            0.957629      0.000547
12     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy',
'rf__max_depth': 20, 'rf__n_estimators': 100}            0.957597      0.000608
13     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini',
'rf__max_depth': 20, 'rf__n_estimators': 50}             0.957559      0.000844
14     {'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini',
'rf__max_depth': 20, 'rf__n_estimators': 50}             0.957461      0.000344
15     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy',
'rf__max_depth': 20, 'rf__n_estimators': 50}             0.957196      0.000585
16     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini',
'rf__max_depth': 10, 'rf__n_estimators': 100}            0.949825      0.000418
17     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'gini',
'rf__max_depth': 10, 'rf__n_estimators': 50}             0.949810      0.000276
18     {'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini',
'rf__max_depth': 10, 'rf__n_estimators': 100}            0.949698      0.000423
19     {'custom_transformer__use_extra_features': False, 'rf__criterion': 'gini',
'rf__max_depth': 10, 'rf__n_estimators': 50}             0.949693      0.000501
20     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy',
'rf__max_depth': 10, 'rf__n_estimators': 100}            0.949271      0.000319
21     {'custom_transformer__use_extra_features': True, 'rf__criterion': 'entropy',
'rf__max_depth': 10, 'rf__n_estimators': 50}             0.949153      0.000421
22     {'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy',
'rf__max_depth': 10, 'rf__n_estimators': 100}            0.948897      0.000655
23     {'custom_transformer__use_extra_features': False, 'rf__criterion': 'entropy',
'rf__max_depth': 10, 'rf__n_estimators': 50}             0.948809      0.000789
24

```

Grid Search Complete in 111.43 seconds.

```
Best Parameters: {'custom_transformer__use_extra_features': False,
'rf__criterion': 'entropy', 'rf__max_depth': None, 'rf__n_estimators': 100}
Validation F1-Score: 0.9615
```

```
[5]: def train_with_history(X_train, y_train, X_val, y_val, best_pipeline):
    print("Preparing data for history tracking...")

    # 1. Extract the best RF parameters and set warm_start
    rf_step = best_pipeline.named_steps['rf']
    rf_params = rf_step.get_params()
    rf_params['warm_start'] = True

    # 2. Extract the preprocessing steps (Transformer + Scaler)
    # This ensures your custom features are included in the chart metrics
    preprocessor = Pipeline(best_pipeline.steps[:-1])
    X_train_prep = preprocessor.transform(X_train)
    X_val_prep = preprocessor.transform(X_val)

    # 3. Initialize the RF with the best parameters (no duplicates)
    rf = RandomForestClassifier(**rf_params)

    train_loss, val_loss = [], []
    train_acc, val_acc = [], []

    # Define the range of trees to track (1 to the best n_estimators)
    max_trees = rf_params.get('n_estimators', 100)
    tree_range = range(1, max_trees + 1, 5)

    print(f"Recording metrics for up to {max_trees} trees...")

    for i in tree_range:
        rf.n_estimators = i
        rf.fit(X_train_prep, y_train)

        # Calculate Log Loss (Cross-Entropy)
        train_loss.append(log_loss(y_train, rf.predict_proba(X_train_prep)))
        val_loss.append(log_loss(y_val, rf.predict_proba(X_val_prep)))

        # Calculate Accuracy
        train_acc.append(accuracy_score(y_train, rf.predict(X_train_prep)))
        val_acc.append(accuracy_score(y_val, rf.predict(X_val_prep)))

    # --- Visualization ---
    plt.figure(figsize=(15, 5))

    # Plot Accuracy
    plt.subplot(1, 2, 1)
```

```

plt.plot(tree_range, train_acc, label='Train Accuracy', marker='.', color="#1f77b4")
plt.plot(tree_range, val_acc, label='Val Accuracy', marker='.', color="#ff7f0e")
plt.title('Accuracy vs. Ensemble Size')
plt.xlabel('Number of Trees')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)

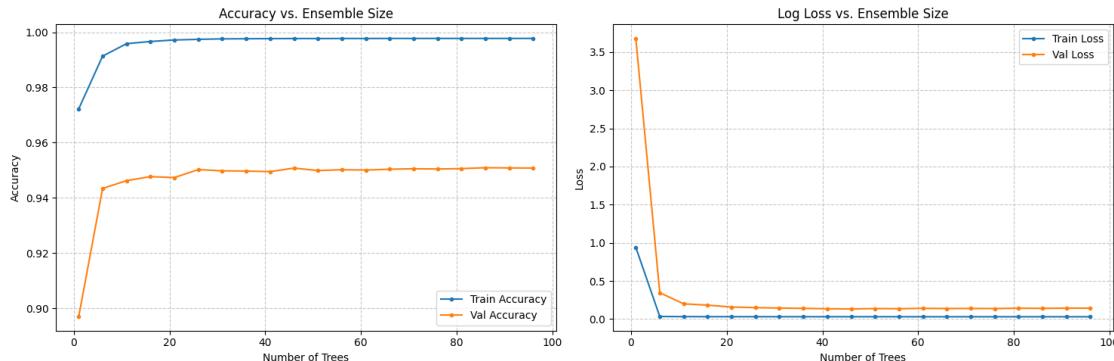
# Plot Loss (Log Loss)
plt.subplot(1, 2, 2)
plt.plot(tree_range, train_loss, label='Train Loss', marker='.', color="#1f77b4")
plt.plot(tree_range, val_loss, label='Val Loss', marker='.', color="#ff7f0e")
plt.title('Log Loss vs. Ensemble Size')
plt.xlabel('Number of Trees')
plt.ylabel('Loss')
plt.legend()
plt.grid(True, linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()

# Run this after your Grid Search
train_with_history(X_train, y_train, X_val, y_val, best_rf_model)

```

Preparing data for history tracking...  
Recording metrics for up to 100 trees...



```
[6]: def evaluate_model(model, X_test, y_test):
    print("\n" + "="*50)
    print(" FINAL PERFORMANCE ON TEST SET ")
```

```

print("=="*50)

y_pred = model.predict(X_test)

print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=['Normal', ↴
'Attack'], digits=4))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Normal', 'Attack'],
            yticklabels=['Normal', 'Attack'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix: Random Forest')
plt.show()

# Execute final evaluation
evaluate_model(best_rf_model, X_test, y_test)

```

=====
FINAL PERFORMANCE ON TEST SET
=====

	precision	recall	f1-score	support
Normal	0.9281	0.9378	0.9329	13950
Attack	0.9647	0.9589	0.9618	24701
accuracy			0.9513	38651
macro avg	0.9464	0.9484	0.9474	38651
weighted avg	0.9515	0.9513	0.9514	38651

Confusion Matrix: Random Forest

