

# task\_1\_svm

February 22, 2026

---

## 0.1 Cell 1: Mount Google Drive & Import Libraries

In this section, we import the necessary Python libraries for data manipulation (`pandas`, `numpy`), visualization (`seaborn`, `matplotlib`), and machine learning (`sklearn`).

**Key Libraries:** \* `LinearSVC`: The primary Support Vector Machine model (optimized for large datasets). \* `CalibratedClassifierCV`: Wraps the SVM to provide probability estimates (required for ROC-AUC). \* `GridSearchCV`: For hyperparameter tuning.

```
[1]: # # # Mount Google Drive
# from google.colab import drive
# drive.mount('/content/drive')

# Core Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import time

# Scikit-Learn Models
from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV

# Scikit-Learn Tools
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import (
    classification_report, confusion_matrix, accuracy_score,
    f1_score, roc_auc_score, recall_score, precision_score
)
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.base import BaseEstimator, TransformerMixin

# Ignore warnings
import warnings
warnings.filterwarnings('ignore')
```

```
print("Libraries imported successfully!")
```

Libraries imported successfully!

## 0.2 Cell 2: Data Preparation & Preprocessing

1. **Data Loading & Merging:** We load both the provided training and testing sets and merge them into a single dataset. This allows us to re-shuffle and split the data consistently.
2. **Data Cleaning:** We drop the `id` and `attack_cat` columns.
  - *Note:* `attack_cat` is removed because it is a granular version of the target label. Including it would cause **data leakage** (giving the model the answer).
3. **Encoding Strategy:**
  - We use **One-Hot Encoding** for categorical variables (`proto`, `service`, `state`).
  - *Justification:* Unlike Random Forest, SVM is a distance-based algorithm. Label Encoding (1, 2, 3) would imply an ordinal relationship (3 > 1) which does not exist, confusing the model. One-Hot Encoding prevents this.
4. **Stratified Splitting:**
  - **Training Set (70%):** For learning patterns.
  - **Validation Set (15%):** For hyperparameter tuning.
  - **Test Set (15%):** For final unseen evaluation.

```
[2]: def load_and_prep_data():  
    # Update paths to match your exact Drive location  
    train_path = '/content/drive/MyDrive/University/CSCI316 - Big Data Mining' +  
    ↪Techniques/Group Assignment/UNSW_NB15_training-set.csv'  
    test_path = '/content/drive/MyDrive/University/CSCI316 - Big Data Mining' +  
    ↪Techniques/Group Assignment/UNSW_NB15_testing-set.csv'  
  
    # Local ENV file path  
    test_path = '/Users/jju/Documents/SIM/Semester 1, 2026/CSCI316 Big Data' +  
    ↪Mining/Assignments/Group Assignment_Database/UNSW_NB15_testing-set.csv'  
    train_path = '/Users/jju/Documents/SIM/Semester 1, 2026/CSCI316 Big Data' +  
    ↪Mining/Assignments/Group Assignment_Database/UNSW_NB15_training-set.csv'  
  
    print("Loading data...")  
    df_train_orig = pd.read_csv(train_path)  
    df_test_orig = pd.read_csv(test_path)  
  
    # 1. Combine Datasets (Matching RF logic)  
    df_full = pd.concat([df_train_orig, df_test_orig], axis=0).  
    ↪reset_index(drop=True)  
  
    # 2. Drop Forbidden Columns (Matching RF logic + Fixes Leakage)  
    drop_cols = ['id', 'label', 'attack_cat']  
    X = df_full.drop(drop_cols, axis=1)
```

```

y = df_full['label']

# 4. Stratified Split: Train (70%) and Temp (30%)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# 5. Split Temp into Validation (15%) and Test (15%)
# 0.5 * 30% = 15%
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)

# Now X_train, X_val, X_test/ y_train, y_val, y_test available

# 5. Encoding Strategy: ONE-HOT (Specific to SVM)
# RF uses LabelEncoding, but SVM needs One-Hot.
categorical_cols = ['proto', 'service', 'state'] # attack_cat removed since
# already dropped
# ONE-HOT Encoding each train/test/validation data
X_train = pd.get_dummies(X_train, columns=categorical_cols, drop_first=True)
X_val = pd.get_dummies(X_val, columns=categorical_cols, drop_first=True)
X_test = pd.get_dummies(X_test, columns=categorical_cols, drop_first=True)

# reindex based on X_train after OHE
X_val = X_val.reindex(columns=X_train.columns, fill_value=0)
X_test = X_test.reindex(columns=X_train.columns, fill_value=0)

print(f"Features used for training {len(X.columns)} total:")
# print(list(X.columns))
print("-" * 50)

print(f"Data Loaded and Split:")
print(f"Train: {X_train.shape}, Val: {X_val.shape}, Test: {X_test.shape}")

return X_train, y_train, X_val, y_val, X_test, y_test

# Execute
X_train, y_train, X_val, y_val, X_test, y_test = load_and_prep_data()

# Check distribution (Optional, matching RF style)
def check_distributions(y_train, y_val, y_test):
    sets = {'Train': y_train, 'Validation': y_val, 'Test': y_test}
    for name, set_data in sets.items():
        counts = set_data.value_counts(normalize=True) * 100
        print(f"{name} Distribution: Normal: {counts[0]:.2f}%, Attack: "
              f"{counts[1]:.2f}%")

```

```

check_distributions(y_train, y_val, y_test)

Loading data...
Features used for training (42 total):
-----
Data Loaded and Split:
Train: (180371, 190), Val: (38651, 190), Test: (38651, 190)
Train Distribution: Normal: 36.09%, Attack: 63.91%
Validation Distribution: Normal: 36.09%, Attack: 63.91%
Test Distribution: Normal: 36.09%, Attack: 63.91%

```

### 0.3 Cell 3: Custom Feature Transformer

As per the assignment requirements, we implement a **custom feature transformer**. We engineer two new features to improve model performance:

1. **Packet Ratio (pkt\_ratio):**
  - Formula:  $(spkts + 1)/(dpkts + 1)$
  - *Logic:* Attacks often have imbalanced packet flows (e.g., DoS floods send many packets but receive few).
2. **TTL Gap (ttl\_gap):**
  - Formula:  $|sttl - ttl|$
  - *Logic:* A large difference in Time-To-Live (TTL) values between source and destination can indicate abnormal network pathing or spoofing.

```
[3]: class CustomFeatureTransformer(BaseEstimator, TransformerMixin):
    def __init__(self, use_extra_features=True):
        self.use_extra_features = use_extra_features

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        X_transformed = X.copy()
        if self.use_extra_features:
            # Matches RF Notebook Logic exactly
            if 'spkts' in X.columns and 'dpkts' in X.columns:
                X_transformed['pkt_ratio'] = (X_transformed['spkts'] + 1) / (X_transformed['dpkts'] + 1)

            if 'sttl' in X.columns and 'dttl' in X.columns:
                X_transformed['ttl_gap'] = abs(X_transformed['sttl'] - X_transformed['dttl'])

        return X_transformed
```

## 0.4 Cell 4: Model Selection & Training (SVM)

We utilize a **Pipeline** to prevent data leakage during cross-validation. The pipeline consists of: 1. **Feature Engineering:** Generating `pkt_ratio` and `ttl_gap`. 2. **Scaling (StandardScaler):** \* *Critical for SVM*: SVM attempts to maximize the margin between classes. If features are on different scales (e.g., `sbytes` is huge, `ttl` is small), the model will be biased. Scaling ensures all features contribute equally. 3. **Model (LinearSVC):** \* We wrap `LinearSVC` in `CalibratedClassifierCV` to enable `predict_proba()`, allowing us to calculate ROC-AUC scores.

**Hyperparameter Tuning:** We use `GridSearchCV` to find the optimal Regularization parameter (`C`) and evaluate whether the extra features improve performance.

```
[4]: def train_and_tune_svm(X_train, y_train, X_val, y_val):
    print("Starting Grid Search for SVM (LinearSVC)...")
    start_time = time.time()

    # Pipeline: FeatureEng -> StandardScale -> SVM
    # Note: Using CalibratedClassifierCV to enable probability prediction for
    ↪AUC
    pipeline = Pipeline([
        ('custom_transformer', CustomFeatureTransformer()),
        ('scaler', StandardScaler()),
        ('svm', CalibratedClassifierCV(
            estimator=LinearSVC(random_state=42, dual=False, max_iter=2000),
            cv=3
        ))
    ])

    # Parameters to tune
    param_grid = {
        'custom_transformer__use_extra_features': [True, False],
        'svm__estimator__C': [0.01, 0.1, 1], # Reduced C values for faster
    ↪convergence
    }

    grid_search = GridSearchCV(
        pipeline,
        param_grid,
        cv=3,
        scoring='f1',
        verbose=1,
        n_jobs=-1
    )

    grid_search.fit(X_train, y_train)

    # --- Printing Results (Matches RF Format) ---
    results_df = pd.DataFrame(grid_search.cv_results_)
```

```

relevant_results = results_df[['params', 'mean_test_score', 'std_test_score', 'rank_test_score']]
relevant_results = relevant_results.sort_values(by='rank_test_score')

print("\n--- All Hyperparameter Combinations (Sorted by F1-Score) ---")
# Clean up params column for display
pd.set_option('display.max_colwidth', None)
print(relevant_results.to_string(index=False))

# Validate
val_predictions = grid_search.best_estimator_.predict(X_val)
val_f1 = f1_score(y_val, val_predictions)

print(f"\nGrid Search Complete in {time.time() - start_time:.2f} seconds.")
print(f"Best Parameters: {grid_search.best_params_}")
print(f"Validation F1-Score: {val_f1:.4f}")

return grid_search.best_estimator_

best_svm_model = train_and_tune_svm(X_train, y_train, X_val, y_val)

```

Starting Grid Search for SVM (LinearSVC)...

Fitting 3 folds for each of 6 candidates, totalling 18 fits

---

All Hyperparameter Combinations (Sorted by F1-Score) ---

mean_test_score	std_test_score	rank_test_score	params
{'custom_transformer__use_extra_features': True, 'svm__estimator__C': 1}			
0.928802	0.000524	1	
{'custom_transformer__use_extra_features': True, 'svm__estimator__C': 0.1}			
0.928364	0.000537	2	
{'custom_transformer__use_extra_features': False, 'svm__estimator__C': 1}			
0.928130	0.000950	3	
{'custom_transformer__use_extra_features': True, 'svm__estimator__C': 0.01}			
0.927131	0.000625	4	
{'custom_transformer__use_extra_features': False, 'svm__estimator__C': 0.1}			
0.927069	0.000541	5	
{'custom_transformer__use_extra_features': False, 'svm__estimator__C': 0.01}			
0.925470	0.000149	6	

Grid Search Complete in 500.35 seconds.

Best Parameters: {'custom\_transformer\_\_use\_extra\_features': True, 'svm\_\_estimator\_\_C': 1}

Validation F1-Score: 0.9286

---

## 0.5 Cell 5: Final Evaluation & Summary

We evaluate the best-performing SVM model on the **held-out Test Set (15%)**.

**Metrics Reported:** \* **Accuracy:** Overall correctness. \* **F1-Score:** Harmonic mean of precision and recall (crucial for imbalanced data). \* **ROC-AUC:** The model's ability to distinguish between Normal and Attack traffic. \* **Confusion Matrix:** To visualize False Positives vs. False Negatives.

```
[5]: def evaluate_model(model, X_test, y_test):
    print("\n" + "=" * 60)
    print("      FINAL EVALUATION ON TEST SET (SVM)")
    print("=" * 60)

    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:, 1]

    # 1. Classification Report
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred, target_names=['Normal', 'Attack'], digits=4))

    # 2. Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
                xticklabels=['Normal', 'Attack'],
                yticklabels=['Normal', 'Attack'])
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title('Confusion Matrix - SVM')
    plt.show()

    # Run Evaluation
    evaluate_model(best_svm_model, X_test, y_test)
```

```
=====
      FINAL EVALUATION ON TEST SET (SVM)
=====

Classification Report:
precision    recall   f1-score   support
Normal       0.9233    0.8106    0.8633    13950
Attack       0.8999    0.9619    0.9299    24701

accuracy          0.9073    38651
macro avg       0.9116    0.8863    0.8966    38651
weighted avg     0.9084    0.9073    0.9059    38651
```

Confusion Matrix - SVM

